

Predicting Bug Resolution on Eclipse Browser

Al-sadh Imadh, Hope Mullins, Gualberto Oliveira

1 August 2024

Summary

We explored a dataset on bug tracking for the Eclipse Browser.

Summary

We explored a dataset on bug tracking for the Eclipse Browser.

Our goal was to analyze the factors that impacted the likelihood of a bug being fixed.

Summary

We explored a dataset on bug tracking for the Eclipse Browser.

Our goal was to analyze the factors that impacted the likelihood of a bug being fixed.

We trained 5 types of models and found that Random Forest gave the best AUC (Area Under the Curve) score.

ROADMAP OF TALK

① Problem

ROADMAP OF TALK

① Problem

② Theory

ROADMAP OF TALK

- ① Problem
- ② Theory
- ③ Data

ROADMAP OF TALK

- ① Problem
- ② Theory
- ③ Data
- ④ Training, Validation, and Testing

ROADMAP OF TALK

- ① Problem
- ② Theory
- ③ Data
- ④ Training, Validation, and Testing
- ⑤ Conclusions

Software Bugs

Bugs are a big problem in software and cause programs to behave in ways that are unexpected or not at all.

Software Bugs

Bugs are a big problem in software and cause programs to behave in ways that are unexpected or not at all.

On average, they are quoted to cost nearly \$60 Billion each year.

Software Bugs

Bugs are a big problem in software and cause programs to behave in ways that are unexpected or not at all.

On average, they are quoted to cost nearly \$60 Billion each year.

Assignments to developers who will work on fixing the bug sometimes choose to not fix bugs due to just how labor intensive they will be.

Bug Prediction

The Eclipse and Mozilla Defect Tracking Dataset contains a plethora of information about bug reports and how they are updated over time. [LPD13]

Bug Prediction

The Eclipse and Mozilla Defect Tracking Dataset contains a plethora of information about bug reports and how they are updated over time. [LPD13]

We will use this data to predict the likelihood of a bug report being fixed or not fixed.

Bug Prediction

The Eclipse and Mozilla Defect Tracking Dataset contains a plethora of information about bug reports and how they are updated over time. [LPD13]

We will use this data to predict the likelihood of a bug report being fixed or not fixed.

To do this, we will employ 5 types of models on an altered dataset and compare their results to draw conclusions.

LOGIT Regression

- 1 This model takes linear combination of the factors and transforms it to return a value $p \in [0, 1]$

- Regression equation:

$$\ln\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 x + \dots + \beta_k x_k$$

LOGIT Regression

- 1 This model takes linear combination of the factors and transforms it to return a value $p \in [0, 1]$
- 2 0 represents "Won't Fix", 1 represents "Fixed"

- Regression equation:

$$\ln\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 x + \dots + \beta_k x_k$$

LOGIT Regression


- 1 This model takes linear combination of the factors and transforms it to return a value $p \in [0, 1]$
 - 2 0 represents "Won't Fix", 1 represents "Fixed"
 - 3 Our optimal threshold to label output as 1 was 0.56
- Regression equation:

$$\ln\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 x + \dots + \beta_k x_k$$

Lasso Regularization

- ① This improves the previous LOGIT model by adding to the loss function L

$$L + \lambda \sum_i |\beta_i|$$




<https://allmodelsarewrong.github.io/lasso.html>

Lasso Regularization

- 1 This improves the previous LOGIT model by adding to the loss function L

$$L + \lambda \sum_i |\beta_i|$$

- 2 Model is encouraged to push (and even set) β_i 's to 0.




<https://allmodelsarewrong.github.io/lasso.html>

Bagging

Also known as Bootstrap Aggregation.

- 1 Reduces variance in a noisy dataset




<https://www.datacamp.com/tutorial/what-bagging-in-machine-learning-a-guide-with-examples>

Bagging

Also known as Bootstrap Aggregation.

- 1 Reduces variance in a noisy dataset
- 2 Initial dataset is bootstrapped and each sample is used to make a decision tree




<https://www.datacamp.com/tutorial/what-bagging-in-machine-learning-a-guide-with-examples>

Bagging

Also known as Bootstrap Aggregation.


- 1 Reduces variance in a noisy dataset
- 2 Initial dataset is bootstrapped and each sample is used to make a decision tree
- 3 Smaller models are averaged/aggregated to make overall model



<https://www.datacamp.com/tutorial/what-bagging-in-machine-learning-a-guide-with-examples>

Random Forest


1 Closely related to Bagging



<https://medium.com/@roiyehe/random-forests-98892261dc49>

Random Forest


- 1 Closely related to Bagging
- 2 Unlike bagging, only a handful of the variables are considered during a node split



<https://medium.com/@roiyehe/random-forests-98892261dc49>

Random Forest


- 1 Closely related to Bagging
- 2 Unlike bagging, only a handful of the variables are considered during a node split
- 3 The variables are chosen randomly during training



<https://medium.com/@roiyehe/random-forests-98892261dc49>

Random Forest

- 1 Closely related to Bagging
- 2 Unlike bagging, only a handful of the variables are considered during a node split
- 3 The variables are chosen randomly during training
- 4 The random selection leads to more independent trees which can help prediction



Boosting

- 1 Weak decision trees are built sequentially



<https://www.geeksforgeeks.org/boosting-in-machine-learning-boosting-and-adaboost/>

Boosting

- 1 Weak decision trees are built sequentially
- 2 Models (Learners) build on top of the previous one, each model weighted on performance



<https://www.geeksforgeeks.org/boosting-in-machine-learning-boosting-and-adaboost/>

Boosting

- 1 Weak decision trees are built sequentially
- 2 Models (Learners) build on top of the previous one, each model weighted on performance
- 3 We used Gradient Boosting, which tries to build learners that are more efficient than the previous



<https://www.geeksforgeeks.org/boosting-in-machine-learning-boosting-and-adaboost/>

Eclipse Bug Tracking Dataset

- The Eclipse and Mozilla Defect Tracking Dataset

Eclipse Bug Tracking Dataset

- The Eclipse and Mozilla Defect Tracking Dataset
- 12 Excel files containing attributes about a bug. Attributes are updated over time

Eclipse Bug Tracking Dataset

- The Eclipse and Mozilla Defect Tracking Dataset
- 12 Excel files containing attributes about a bug. Attributes are updated over time
- Only bugs labeled "Fixed" or "Won't Fix" are kept. Duplicate bugs removed. Final dataset has size of 90,797

Eclipse Bug Tracking Dataset

- The Eclipse and Mozilla Defect Tracking Dataset
- 12 Excel files containing attributes about a bug. Attributes are updated over time
- Only bugs labeled "Fixed" or "Won't Fix" are kept. Duplicate bugs removed. Final dataset has size of 90,797
- New variables created using information in dataset

Eclipse Bug Tracking Dataset

- The Eclipse and Mozilla Defect Tracking Dataset
- 12 Excel files containing attributes about a bug. Attributes are updated over time
- Only bugs labeled "Fixed" or "Won't Fix" are kept. Duplicate bugs removed. Final dataset has size of 90,797
- New variables created using information in dataset
- Excel and SQLite3 were used to cleanup the data and create new variables

Data Structure

id	curr_res	reporter	stat_upd	num_intrst	op_sys	component	prod	severity
287149	FIXED	17941	2	6	Linux-GTK	SWT	Platform	normal
89374	FIXED	57	5	1	Windows 2000	UI	Platform	normal
89378	FIXED	8126	3	5	Windows XP	SWT	Platform	normal

Sample of dataset. Some new variables include:

Data Structure

id	curr_res	reporter	stat_upd	num_intrst	op_sys	component	prod	severity
287149	FIXED	17941	2	6	Linux-GTK	SWT	Platform	normal
89374	FIXED	57	5	1	Windows 2000	UI	Platform	normal
89378	FIXED	8126	3	5	Windows XP	SWT	Platform	normal

Sample of dataset. Some new variables include:

- Total time bug report is open
- Max priority/severity
- Number of reassignments/status changes
- Success rate of initial assignee


Variable Descriptions (1)

- id: Unique identifier for bug
- curr_res: Whether a bug is fixed or won't be fixed
- reporter: Unique identifier for the bug reporter
- stat_upd: Number of times the bug status was updated
- num_inrst: Number of emails interested in the bug
- op_sys: Operating system that bug affects. Most recent update value used
- prod: Software/product that the bug pertains to
- component: Subsystem of product that the bug affects
- severity: Highest severity given to the bug

Variable Description (2)

- version: The version of the product that the bug affects
- times_assigned: Number of times the bug was reassigned
- succ_rate: The success rate of the initial assignee
- res_upd: Number of times the resolution of a bug was changed
- res_time: Time until the bug was resolved
- reporter_report_cnt: How many bugs the reporter has reported in the dataset
- desc_length: Sum of lengths of the descriptions on the bug
- prio: Highest priority value assigned to the bug

Model: LOGIT Regression (with poly)

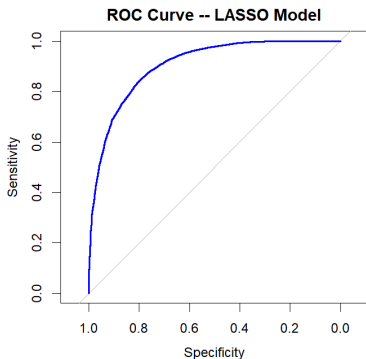


Confusion Matrix

class	Wont Fix	Fixed
Wont Fix	2848	777
Fixed	3070	29624

Figure: AUC: 0.9048

Model: LASSO Regularization (with poly)

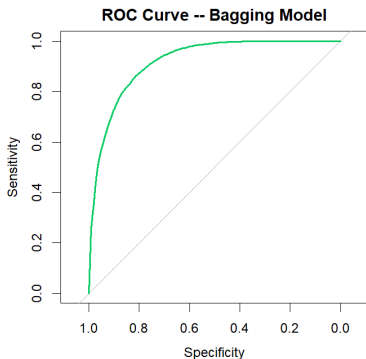


Confusion Matrix

class	Wont Fix	Fixed
Wont Fix	3161	826
Fixed	2757	29575

Figure: AUC: 0.9286

Model: Bagging




Confusion Matrix

class	Wont Fix	Fixed
Wont Fix	3425	521
Fixed	2493	29880

Figure: AUC: 0.9186

Model: Random Forest




Confusion Matrix

class	Wont Fix	Fixed
Wont Fix	3431	438
Fixed	2487	29963

Figure: AUC: 0.9286

Model: Boosting



Confusion Matrix

class	Wont Fix	Fixed
Wont Fix	3349	916
Fixed	2569	29485

Figure: AUC: 0.9122

Important Variables

- 1 Based on the results of the LOGIT regression and the subsequent LASSO regularization

Important Variables

- 1 Based on the results of the LOGIT regression and the subsequent LASSO regularization
- 2 Almost every variable we examined was significant

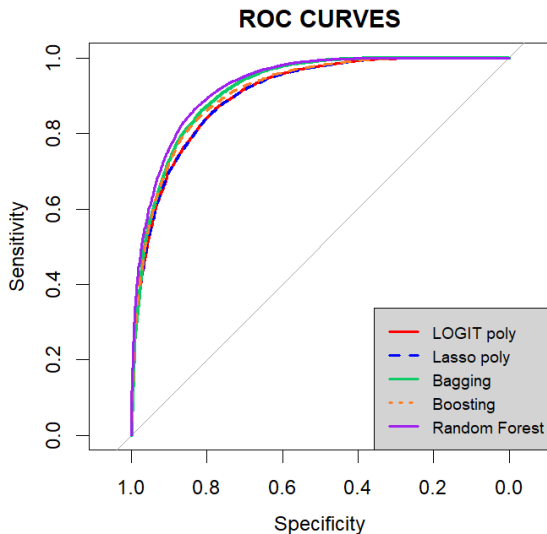
Important Variables

- 1 Based on the results of the LOGIT regression and the subsequent LASSO regularization
- 2 Almost every variable we examined was significant
- 3 The only variable that was found to not be significant was "Component"

Summary of LOGIT variables (no poly)

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)	-2.4051	0.1819	-13.22	0.0000
reporter	-0.0000	0.0000	-19.40	0.0000
stat_upd	0.5816	0.0208	27.90	0.0000
num_inrst	0.1420	0.0095	14.99	0.0000
op_sys	0.0039	0.0012	3.25	0.0012
component	-0.0007	0.0016	-0.40	0.6888
prod	-0.1186	0.0179	-6.62	0.0000
severity	0.2028	0.0113	17.90	0.0000
version	0.0382	0.0025	15.53	0.0000
times_assigned	0.2017	0.0179	11.28	0.0000
succ_rate	5.1959	0.1233	42.15	0.0000
res_upd	-0.7482	0.0322	-23.26	0.0000
res_time	-0.0000	0.0000	-46.34	0.0000
reporter_report_cnt	-0.0002	0.0000	-15.26	0.0000
desc_length	-0.0022	0.0003	-6.46	0.0000
prio	-0.1457	0.0151	-9.65	0.0000

ROC Curve Comparison



Selection of Best Model

- 1 All models but bagging were also tested using `poly()` to give 135 variables (including polynomial combinations and interaction effects), with LOGIT and LASSO seeing the most notable improvements

Selection of Best Model

- 1 All models but bagging were also tested using `poly()` to give 135 variables (including polynomial combinations and interaction effects), with LOGIT and LASSO seeing the most notable improvements
- 2 Boosting also improved, but not enough to outperform our Random Forest model

Selection of Best Model

- ① All models but bagging were also tested using `poly()` to give 135 variables (including polynomial combinations and interaction effects), with LOGIT and LASSO seeing the most notable improvements
- ② Boosting also improved, but not enough to outperform our Random Forest model
- ③ We used the Area Under the Curve to determine our best model. Without `poly()`, Boosting, Random Forest, and Bagging all returned similarly high AUC scores

Selection of Best Model

- 1 All models but bagging were also tested using `poly()` to give 135 variables (including polynomial combinations and interaction effects), with LOGIT and LASSO seeing the most notable improvements
- 2 Boosting also improved, but not enough to outperform our Random Forest model
- 3 We used the Area Under the Curve to determine our best model. Without `poly()`, Boosting, Random Forest, and Bagging all returned similarly high AUC scores
- 4 Using our training/testing set, Random Forest gave the highest AUC, i.e. the best testing performance, regardless of whether we utilized the `poly()` function or not.

Final Notes

- 1 Each model was a good predictor of how a bug would be resolved

Final Notes

- ① Each model was a good predictor of how a bug would be resolved
- ② The confusion matrix shows that every model struggles to identify when a bug will not be fixed correctly (considering that the majority of the data consisted of "FIXED" bugs, this makes sense)

Final Notes

- ① Each model was a good predictor of how a bug would be resolved
- ② The confusion matrix shows that every model struggles to identify when a bug will not be fixed correctly (considering that the majority of the data consisted of "FIXED" bugs, this makes sense)
- ③ In line with the AUC scores, Random Forest identifies both categories the best with Bagging at a close second

Final Notes

- 1 Each model was a good predictor of how a bug would be resolved
- 2 The confusion matrix shows that every model struggles to identify when a bug will not be fixed correctly (considering that the majority of the data consisted of "FIXED" bugs, this makes sense)
- 3 In line with the AUC scores, Random Forest identifies both categories the best with Bagging at a close second
- 4 Boosting was expected to outperform every other model but ranked below both Bagging and Random Forest

Final Notes

- 1 Each model was a good predictor of how a bug would be resolved
- 2 The confusion matrix shows that every model struggles to identify when a bug will not be fixed correctly (considering that the majority of the data consisted of "FIXED" bugs, this makes sense)
- 3 In line with the AUC scores, Random Forest identifies both categories the best with Bagging at a close second
- 4 Boosting was expected to outperform every other model but ranked below both Bagging and Random Forest
- 5 Rankings could change with different val/train splits

Citation



Ahmed Lamkanfi, Javier Perez, and Serge Demeyer, *The eclipse and mozilla defect tracking dataset: a genuine dataset for mining bug information*, MSR '13: Proceedings of the 10th Working Conference on Mining Software Repositories, May 18–19, 2013. San Francisco, California, USA, 2013.